

DATA SCIENCE DISCUSSION BRIEF

Fluency Tracker Measurement Model

A practical view of how Glean telemetry becomes aggregate, defensible AI fluency evidence without raw content capture, individual scoring, or unsupported productivity claims.

Recommended framing

Glean shows what happened. Fluency Tracker decides what can be interpreted from it. Data Science validates whether those interpretations are credible enough to use in measurement.

AUDIENCE

Juhi Singh, Data Science Lead

USE CASE

AI fluency measurement model review

UNIT OF ANALYSIS

Workflow execution, aggregated to org/window

BOUNDARY

Metadata and event sequences only

The Measurement Model

Aggregate only

The model separates source telemetry, derived behavioral evidence, validation, and claim governance. This is the core picture to socialize with Data Science.

1. GLEAN COLLECTS

Structured product events

- Assistant and chat sessions
- Search and retrieval activity
- Workflow, agent, action, and MCP runs

- Security, policy, admin, and product snapshots where available

V

2. FLUENCY TRACKER DERIVES

Behavioral proxies

- Signal readiness and joinability
- Canonical event normalization
- Trace reconstruction across starts, steps, checks, errors, retries, and completions
- Feature families mapped to fluency constructs

V

3. DATA SCIENCE VALIDATES

Construct credibility

- Does each proxy mean what we claim?
- Where are missingness, bias, and drift risks?
- Which thresholds are stable enough for reporting?
- Which claims should stay suppressed?

V

4. OUTPUTS ARE GOVERNED

Claim-safe evidence

- Aggregate workflow behavior patterns
- Enablement and governance priorities
- Suppression when evidence is thin, ambiguous, non-joinable, or privacy sensitive

How Much Is Glean vs. Fluency Tracker?

Glean is the source of many observed product events. Fluency Tracker is the measurement and governance layer that turns those events into validated behavioral evidence.

LAYER	WHAT EXISTS OR IS COLLECTED	WHO OWNS THE INTERPRETATION	MEASUREMENT POSTURE
Source telemetry	Glean product logs: assistant/chat, search, retrieval, workflow runs, agent steps, actions, MCP, policy/security events, and product snapshots where enabled.	Glean owns source instrumentation and access contracts.	Observed activity, not fluency by itself.
Readiness map	Availability, scrub status, join keys, completeness, time semantics, and org-window coverage by signal family.	Fluency Tracker plus Glean data owners.	Determines which signals are usable for a given customer/window.
Behavioral derivation	Ordered workflow traces and passive behavior proxies derived from safe metadata and event sequences.	Fluency Tracker.	Feature engineering layer, not a productivity model.
Construct validation	Proxy confidence, coverage, drift, bias checks, threshold review, and claim approval.	Data Science.	Decides what is credible enough to present.

Passive Behaviors Tracked

No raw content required

These are metadata-derived behavior proxies. They are not person-level judgments, output-quality scores, or causal productivity claims.

Adoption and Exposure

First observed tool use, exploratory feature trials, template reuse, instrumented source coverage.

Work Shaping

Prompt strategy shifts, early constraints, acceptance criteria, scope narrowing, task decomposition.

Trust Calibration

Verification starts, verification latency buckets, counterfactual checks, policy checks, post-run review.

Friction and Iteration

Rapid abandonment, dense retry loops, undo churn, latency buckets, repeat failure loops.

Recovery Behavior

Contradiction detection, conflict resolution, error-to-retry-to-success paths, agent redirects.

Diffusion and Reuse

Artifact forking, peer artifact engagement, template reuse, reusable output creation.

Risk and Oversight

Safe-path escalation, sensitive mode, review before agent execution, stop interventions, permission elevation attempts.

Capability Growth

Movement from basic assistance to reusable workflows, Skills, agentic execution, and governed automation patterns.

What Can and Cannot Be Claimed

The brief is strongest when the claim boundary is explicit. Fluency Tracker should fail closed when evidence is not ready.

Claim-safe with caveats

Observed aggregate workflow behavior

Safe language: covered Glean telemetry can support directional evidence of usage quality, verification behavior, friction, recovery, reuse, oversight, and workflow-level adoption patterns.

Suppress until validated

Unsupported performance or ROI claims

Avoid: individual fluency scoring, team rankings, productivity attribution, causal impact, raw content quality, or ROI translation unless approved assumptions and complete coverage are in place.

Questions For Data Science

This is the review agenda for Juhi: validate the construct map before broadening reporting or executive claims.

QUESTION	WHY IT MATTERS	DEFAULT ANSWER
Is workflow execution the right unit?	Prevents the model from becoming person-level monitoring.	Yes for V1: aggregate to workflow, function, org, and window.

QUESTION	WHY IT MATTERS	DEFAULT ANSWER
Are these credible proxies?	Determines whether adoption quality and fluency claims are defensible.	Validate feature families first: exposure, shaping, calibration, friction, recovery, reuse, oversight.
How should missingness behave?	Missing telemetry can create false confidence.	Fail closed through readiness, full-state coverage, minimum-signal, ambiguity, and cohort-size gates.
What should be reported first?	Sequencing prevents overclaiming.	Start with aggregate behavior coverage and workflow pattern mix. Delay ROI, causality, and productivity language.

What Fluency Tracker Adds

Glean captures structured activity. Fluency Tracker turns that activity into aggregate, validated, claim-safe behavioral evidence.

Data boundary

This model uses metadata and event sequences only. It does not require:

- Employee rows, names, emails, or direct user identifiers.
- Raw prompts, responses, documents, screenshots, or transcripts.
- Person-level rankings, output-quality scoring, or productivity attribution.

Status: Draft measurement brief for Data Science review. Source anchors: docs/fluencytracr-how-it-works-and-behavior-tracking.md , docs/behaviors/V0_Behaviors_and_Formulas.md , and docs/contracts/reportability/operating-handoff.md .